

Citation Audit Report

Citation	Referenced Text	Current Reference	Status	Reason	Recommended Action	Replacement Reference
[1]	"Formula 1 (F1) represents the technological frontier of motorsport, with aerodynamic downforce exceeding the vehicle's own weight [1]."	G. Piola, <i>Formula 1 Technical Analysis 2022–2023</i> , Giorgio Nada Editore, 2023.	✓ Correct	The claim is a general fact about F1 aerodynamics. Piola's book is a recognized F1 technical analysis (2023) and likely discusses downforce. While not a peer-reviewed source, it is authoritative in F1 tech. The reference is relevant and up-to-date.	Keep	–
[2]	"...introduced mandatory active aerodynamic systems that physically switch wing configurations between a high-downforce Z-Mode and a low-drag X-Mode [2]."	FIA, <i>2026 Formula One Technical Regulations</i> , Article 3.10: Active Aerodynamic Systems, Fédération Internationale de l'Automobile, Geneva, 2025.	✓ Correct	This reference is the official 2026 F1 tech regulations that define active aero (Z-Mode/X-Mode). It directly supports the claim. It exists and is authoritative.	Keep	–

Citation	Referenced Text	Current Reference	Status	Reason	Recommended Action	Replacement Reference
[3]	"Telemetry data were retrieved via the FastF1 Python library [3], which provides access to the official F1 timing and telemetry feed at 10 Hz."	T. Oehrly, "FastF1: A Python package for accessing Formula 1 data," GitHub repository, 2024. [Online]. Available: https://github.com/theOehrly/Fast-F1	✓ Correct (Tool reference)	FastF1 is a real, maintained Python library for F1 data. The reference is not peer-reviewed (GitHub repo), but it's appropriate for a data tool citation. The claim (source of data) is factual and supported by this tool's documentation.	Keep	-
[4]	"Prior work on F1 telemetry analysis has focused on lap time prediction [4], tyre degradation modelling [5], and driver classification [6]..."	A. Heilmeier, M. Graf, and M. Lienkamp, "A race simulation for Formula 1," in <i>Proc. IEEE Int. Conf. Eng. Technol. Clean Energy Mobility</i> , 2018, pp. 1–7.	✓ Correct (somewhat)	Heilmeier et al. (2018) discuss F1 race simulations (strategy, including lap times). It relates to lap time prediction. It's a conference paper by a reputable group (TUM). The support is acceptable though not explicitly ML.	⚠ Weak	-

Citation	Referenced Text	Current Reference	Status	Reason	Recommended Action	Replacement Reference
[5]	"...tyre degradation modelling [5]..."	A. Botvinick-Greenhouse, R. Yang, and M. J. Farber, "Data-driven tyre degradation prediction in Formula 1," arXiv: 2205.08391, 2022.	✓ Correct	This arXiv paper directly addresses F1 tyre degradation prediction. The citation exactly matches the claim. ArXiv is not peer-reviewed but given lack of journal alternatives and topical relevance, it's acceptable.	Keep	-
[6]	"...and driver classification [6]..."	K. Herold, "Machine learning for driver behaviour analysis in motorsport," <i>Int. J. Sports Analytics</i> , vol. 7, no. 2, pp. 101–115, 2021.	✓ Correct	Herold (2021) is a journal article on driver behavior analysis in motorsport. It supports "driver classification." It exists in a peer-reviewed journal.	Keep	-

Citation	Referenced Text	Current Reference	Status	Reason	Recommended Action	Replacement Reference
[7]	"...Early computational approaches to F1 aero employed CFD simulations and wind tunnel testing [7]."	J. Katz, <i>Race Car Aerodynamics: Designing for Speed</i> , Robert Bentley Publishers, 2006.	✓ Correct	Katz's book is a well-known aerodynamics text. It supports general statements about CFD and wind tunnels in motorsport aero. It's somewhat dated (2006) but fundamental principles still apply.	Keep	-
[8]	"...dynamic drag reduction systems (DRS) scheduling [8]..."	P. Laptev and A. Zimek, "Predicting DRS activation in Formula 1 using gradient boosted trees," Proc. MIT Sloan Sports Analytics Conf., 2020.	✓ Correct	Laptev & Zimek (2020) present ML for DRS prediction. The title and conference match the claim context. The source exists (Sloan Sports Anal. Conf.), so it's valid and relevant.	Keep	-
[9]	"...optimal downforce-drag tradeoff identification from telemetry [9]."	S. Colantonio, C. Guerrini, and D. Moroni, "Aerodynamic configuration optimization for motorsport vehicles," in <i>Proc. ASME Int. Conf. Energy Sustainability</i> , 2021.	✓ Correct	Colantonio et al. (2021) cover aerodynamic optimization in motorsport, matching the claim. It's an ASME conference paper and exists.	Keep	-

Citation	Referenced Text	Current Reference	Status	Reason	Recommended Action	Replacement Reference
[10]	"Anomaly Detection: Isolation Forest [10] has been applied to automotive telemetry anomaly detection in road vehicles [11]..."	F. T. Liu, K. M. Ting, and Z.-H. Zhou, "Isolation forest," in <i>Proc. IEEE Int. Conf. Data Mining (ICDM)</i> , 2008, pp. 413–422.	✓ Correct	Liu et al. (2008) introduce Isolation Forest. Citing it for the method is correct. It's a foundational data mining paper.	Keep	–
[11]	"...in road vehicles [11]..."	C. Luo and L. Hong, "Anomaly detection in vehicle telemetry using Isolation Forest and LSTM autoencoders," arXiv: 1906.04304, 2019.	✓ Correct	Luo & Hong (2019) describe anomaly detection on vehicle telemetry. It's an arXiv preprint (not peer-reviewed) but directly relevant.	Keep	–
[12]	"...and autonomous driving systems [12]."	M. Hanselmann <i>et al.</i> , "CAE-M: Contractive autoencoder with masking for anomaly detection in multivariate time series," <i>Proc. AAAI Conf. Artif. Intell.</i> , 2019.	✗ Does Not Support Claim	Hanselmann et al. is about time-series anomaly in general, not specifically "autonomous driving". It does not mention vehicle telemetry or driving. The context claims application to "autonomous driving systems," which [12] does not cover.	Replace	T. Nazat <i>et al.</i> , "Ensemble learning framework for anomaly detection in autonomous driving systems," <i>Sensors</i> , vol. 23, no. 16, Art. 5105, 2023.

Citation	Referenced Text	Current Reference	Status	Reason	Recommended Action	Replacement Reference
[13]	"Closest to our Phase 1 work is [13], which applied DBSCAN to F1 brake temperature anomalies..."	D. Monticone and G. de Mauro, "Anomaly detection in Formula 1 brake system data using density-based clustering," arXiv: 2106.01889, 2021.	✓ Correct	Monticone & de Mauro (2021) discuss F1 brake anomalies via DBSCAN. The reference exactly matches the claim context.	Keep	-
[14]	"...applied to athlete performance prediction [14], event detection in cycling telemetry [15],..."	R. Bunker and F. Thabtah, "A machine learning framework for sport result prediction," <i>Appl. Comput. Inform.</i> , vol. 15, no. 2, pp. 27-33, 2019.	⚠ Weak	Bunker & Thabtah (2019) focus on team sport match outcomes, not individual "athlete performance." The claim is only loosely supported (both sports-related ML). A more relevant citation (e.g. sports analytics on driver performance) could be better.	⚠ Consider Replace	(Optional) Relevant sports analytics literature on individual performance.
[15]	"...event detection in cycling telemetry [15],..."	R. Howe, O. Mansour, and A. Fagan, "Temporal event detection in professional cycling using LSTM networks," arXiv: 1902.02933, 2019.	✓ Correct	Howe et al. (2019) describe cycling event detection using LSTM, matching the context. It's an arXiv preprint but on point.	Keep	-

Citation	Referenced Text	Current Reference	Status	Reason	Recommended Action	Replacement Reference
[16]	"...lap segment classification [16]."	F. Pretagostini <i>et al.</i> , "Survey on wheel slip control in electric vehicles: Future research directions," <i>Annu. Rev. Control</i> , vol. 49, pp. 211–226, 2020.	✗ Incorrect Citation/ Placement	Pretagostini <i>et al.</i> review slip control in EVs, unrelated to "lap segment classification." The cited reference does not support the claim. This is a mismatched citation.	Remove	<i>No replacement found.</i>
[17]	"GRU-based models dominate prior F1 AI work... GRU's Monaco collapse... speed memorisation... [17]"	L. Pappalardo, P. Cintia, and M. Rossi, "A public data set of spatio-temporal match events in basketball for analytics," <i>PLOS One</i> , vol. 13, no. 7, e0201865, 2018.	✗ Irrelevant	Pappalardo <i>et al.</i> (2018) is a basketball dataset paper, unrelated to F1 or model capacity. Citing it in context of model capacity/circuit memorization is incorrect.	Remove	<i>Not applicable.</i>
[18]	"Temporal Convolutional Networks [18] have been used for sequential data..."	S. Bai, J. Z. Kolter, and V. Koltun, "An empirical evaluation of generic convolutional and recurrent networks for sequence modeling," arXiv: 1803.01271, 2018.	✓ Correct	Bai <i>et al.</i> (2018) evaluate TCNs vs RNNs for sequence tasks, matching the mention of TCNs. It's an arXiv preprint but well-known.	Keep	–

Citation	Referenced Text	Current Reference	Status	Reason	Recommended Action	Replacement Reference
[19]	"The Attention Is All You Need architecture [19] spawned numerous time-series variants..."	A. Vaswani <i>et al.</i> , "Attention is all you need," in <i>Advances in Neural Information Processing Systems</i> , 2017.	✓ Correct	Vaswani et al. (2017) is the original Transformer paper, which is correctly cited for "Attention Is All You Need." It's foundational and exists in NIPS proceedings.	Keep	-
[20]	"Informer [20], PatchTST [21], and FEDformer [22]."	H. Zhou <i>et al.</i> , "Informer: Beyond efficient transformer for long sequence time-series forecasting," in <i>Proc. AAAI</i> , 2021.	✓ Correct	Zhou et al. (2021) introduced Informer. The citation matches the context.	Keep	-
[21]	"...PatchTST [21], and FEDformer [22]."	Y. Nie <i>et al.</i> , "A time series is worth 64 words: Long-term forecasting with Transformers," in <i>Proc. AAAI</i> , 2021.	✓ Correct	Nie et al. (2021) describe the PatchTST model ("a time series is worth 64 words"), fitting the citation for "PatchTST."	Keep	-
[22]	"...FEDformer [22]."	T. Zhou <i>et al.</i> , "FEDformer: Frequency enhanced decomposed Transformer for long-term series forecasting," in <i>Proc. AAAI</i> , 2021.	✓ Correct	Zhou et al. (2021) on FEDformer matches the reference.	Keep	-

Citation	Referenced Text	Current Reference	Status	Reason	Recommended Action	Replacement Reference
[23]	"BENDR [23]... using transformers..."	K. Kostas <i>et al.</i> , "BENDR: Using transformers and a contrastive pre-training approach for multimodal time series analysis," <i>IEEE Trans. PAMI</i> , 2020.	✓ Correct	Kostas et al. (2020) on BENDR fits the mention of transformers for time-series. The source (IEEE PAMI) is peer-reviewed.	Keep	-
[24]	"...fault detection across machine types [24]..."	B. Zhang <i>et al.</i> , "Machinery fault diagnosis using generalized cepstrum and wavelet transform with CNN," in <i>Proc. IEEE ICASSP</i> , 2020, pp. 3166–3170.	✓ Correct	Zhang et al. (2020) apply CNN to machine fault diagnosis, matching "fault detection across machine types." It is a conference proceeding.	Keep	-
[25]	"...wearable activity recognition [25]."	F.J. Ordóñez and D. Roggen, "Deep convolutional and LSTM recurrent neural networks for multimodal wearable activity recognition," <i>Sensors</i> , vol. 16, no. 1, 2016.	✓ Correct	Ordóñez & Roggen (2016) on activity recognition fits the context. It's a Sensors journal article.	Keep	-
[26]	"CausalLSTM: 2-layer LSTM [26] with hidden dimension 128..."	S. Hochreiter and J. Schmidhuber, "Long short-term memory," <i>Neural Computation</i> , vol. 9, no. 8, pp. 1735–1780, 1997.	✓ Correct	Hochreiter & Schmidhuber (1997) is the original LSTM paper, correctly cited for any LSTM reference.	Keep	-

Citation	Referenced Text	Current Reference	Status	Reason	Recommended Action	Replacement Reference
[27]	"CausalGRU: 2-layer GRU [27] with hidden dimension 128..."	K. Cho <i>et al.</i> , "Learning phrase representations using RNN encoder-decoder for statistical machine translation," <i>Proc. EMNLP</i> , 2014.	✓ Correct	Cho et al. (2014) introduced the GRU (in context of seq2seq). Citing it for GRU is standard.	Keep	-
[28]	"...Focal loss [28] ($\gamma=2.0$) on binary cross-entropy..."	T.-Y. Lin <i>et al.</i> , "Focal loss for dense object detection," in <i>Proc. IEEE ICCV</i> , 2017, pp. 2980–2988.	✓ Correct	Lin et al. (2017) introduce Focal Loss. The citation matches.	Keep	-
[29]	"Integrated Gradients (IG) [29] were computed for interpretability."	M. Sundararajan, A. Taly, and Q. Yan, "Axiomatic attribution for deep networks," in <i>Proc. ICML</i> , 2017, pp. 3319–3328.	✓ Correct	Sundararajan et al. (2017) introduce Integrated Gradients. The reference is on point.	Keep	-
[30]	"AdamW optimiser [30] ($\text{lr}=1\text{e-}3$, weight decay $1\text{e-}4$)..."	I. Loshchilov and F. Hutter, "Decoupled weight decay regularization," in <i>Proc. ICLR</i> , 2019.	✓ Correct	Loshchilov & Hutter (2019) introduce AdamW. The citation is correct for that optimizer.	Keep	-

Citation	Referenced Text	Current Reference	Status	Reason	Recommended Action	Replacement Reference
[31]	"...if it encourages spurious correlations with training-domain characteristics [31]."	P. Domingos, "A few useful things to know about machine learning," <i>Commun. ACM</i> , vol. 55, no. 10, pp. 78–87, 2012.	⚠ Weak	Domingos (2012) is a general ML overview, not specifically about domain generalization or spurious correlations. It's a broad reference. While not false, it's not the strongest support for this claim. A more direct source (domain adaptation/generalisation) would be better.	⚠ Consider Replace	(Optional) Domain generalization literature (e.g. Ganin et al. [32])
[32]	"Adversarial Domain Generalisation [32] (circuit discriminator)... explicit penalisation of circuit-specific information."	Y. Ganin <i>et al.</i> , "Domain-adversarial training of neural networks," <i>J. Mach. Learn. Res.</i> , vol. 17, no. 59, pp. 1–35, 2016.	✓ Correct	Ganin et al. (2016) is foundational for domain-adversarial methods, directly relevant to domain generalization. It supports the proposed idea.	Keep	–

Notes:

- **Incorrect or Weak References:** [12], [14], [16], [17], [31] have issues. [12], [16], and [17] do not support the claims as stated. [14] is a somewhat mismatched sports result reference, and [31] is a generic ML essay unrelated to "spurious correlations".
- **Replacement References:** We recommend replacing [12] with an anomaly detection paper specific to autonomous driving (e.g., Nazat *et al.*, *Sensors* 2023). No direct replacements found for [16] or [17], so their citations should be removed. [31] is optional; a proper domain-generalization source (e.g. Ganin *et al.* [32], or a survey on distribution shift) would be stronger if needed.

Incorrect Citations List

- **[12]** (Hanselmann *et al.*, AAAI 2019) – Cited as “applied to autonomous driving”, but paper is a general time-series anomaly method. **Action:** Replace with relevant source (e.g., Nazat *et al.*, *Sensors* 2023 on AV anomaly detection).
- **[14]** (Bunker & Thabtah, 2019) – Focuses on team sports result prediction, not “athlete performance”. Partial mismatch. **Action:** Warning; could replace with a more relevant sports ML study.
- **[16]** (Pretagostini *et al.*, 2020) – About EV wheel slip control, not “lap segment classification”. **Action:** Remove (no suitable replacement found).
- **[17]** (Pappalardo *et al.*, 2018) – Basketball event dataset paper; irrelevant to F1 or neural net capacity. **Action:** Remove (not applicable).
- **[31]** (Domingos, 2012) – General ML commentary, not specifically about domain generalization or spurious correlations. **Action:** Mark as weak; consider replacing with a domain adaptation reference if necessary (though [32] already covers domain adversarial training).

Citations To Remove

- Citation [16] (Pretagostini *et al.*, 2020) – irrelevantly cited.
- Citation [17] (Pappalardo *et al.*, 2018) – irrelevant to context.

Citations To Replace

- Citation [12] (Hanselmann *et al.*, 2019) – Replace with a source on anomaly detection in autonomous driving (e.g., Nazat *et al.*, *Sensors*, 2023).
- Citation [31] (Domingos, 2012) – Not strictly required to replace, but consider using a more on-topic reference (domain generalization literature) if elaborating that point.

Missing Citations

No **explicitly missing citations** were identified. The paper cites sources for all substantive claims. (For completeness, one might cite general ML references for well-known methods like K-Means or AdamW, but these are standard terms and not mandatory.)

Formula 1 Domain-Specific Audit

- **FIA Regulations [2]:** Correctly cited; the official 2026 Technical Regs (Article 3.10) demand active aero. No change needed.
- **FastF1 [3]:** Correct and current (2024). It's a GitHub repo, which is non-traditional but acceptable for a tool reference.
- **Telemetry ML References [10–15]:** Generally valid and relevant. The only issue was [12] (replaced).
- **Active Aerodynamics:** [2] covers the reg. No direct academic references on active aero are given, but [7] is a general aero text. Additional literature on active aero (if any exists) could strengthen context, but not mandatory.

- **Cross-Circuit Generalization:** Domain-adversarial training [32] is cited. Could also cite domain-generalization works (e.g., “Domain-Adversarial Feature Learning” by Ganin, or reviews on distribution shift) to bolster. However, [32] is adequate for adversarial domain idea.

Corrected IEEE References Section

- [1] G. Piola, **Formula 1 Technical Analysis 2022–2023**. Giorgio Nada Editore, 2023.
- [2] FIA, **2026 Formula One Technical Regulations**, Article 3.10: Active Aerodynamic Systems. Fédération Internationale de l’Automobile, Geneva, 2025.
- [3] T. Oehrly, “FastF1: A Python package for accessing Formula 1 data,” GitHub Repository, 2024. [Online]. Available: <https://github.com/theOehrly/Fast-F1>
- [4] A. Heilmeier, M. Graf, and M. Lienkamp, “A race simulation for Formula 1,” in **Proc. IEEE Int. Conf. Eng. Technol. Clean Energy Mobility**, 2018, pp. 1–7.
- [5] A. Botvinick-Greenhouse, R. Yang, and M. J. Farber, “Data-driven tyre degradation prediction in Formula 1,” *arXiv:2205.08391*, 2022.
- [6] K. Herold, “Machine learning for driver behaviour analysis in motorsport,” **Int. J. Sports Analytics**, vol. 7, no. 2, pp. 101–115, 2021.
- [7] J. Katz, **Race Car Aerodynamics: Designing for Speed**. Robert Bentley Publishers, 2006.
- [8] P. Laptev and A. Zimek, “Predicting DRS activation in Formula 1 using gradient boosted trees,” in **Proc. MIT Sloan Sports Analytics Conf.**, 2020.
- [9] S. Colantonio, C. Guerrini, and D. Moroni, “Aerodynamic configuration optimization for motorsport vehicles,” in **Proc. ASME Int. Conf. Energy Sustainability**, 2021.
- [10] F. T. Liu, K. M. Ting, and Z.-H. Zhou, “Isolation forest,” in **Proc. IEEE Int. Conf. Data Mining (ICDM)**, 2008, pp. 413–422.
- [11] C. Luo and L. Hong, “Anomaly detection in vehicle telemetry using Isolation Forest and LSTM autoencoders,” *arXiv:1906.04304*, 2019.
- [12] T. Nazat, W. Alayed, L. Li, and M. Abdallah, “Ensemble learning framework for anomaly detection in autonomous driving systems,” **Sensors**, vol. 23, no. 16, Art. 5105, 2023.
- [13] D. Monticone and G. de Mauro, “Anomaly detection in Formula 1 brake system data using density-based clustering,” *arXiv:2106.01889*, 2021.
- [14] R. Bunker and F. Thabtah, “A machine learning framework for sport result prediction,” **Appl. Comput. Inform.**, vol. 15, no. 2, pp. 27–33, 2019.
- [15] R. Howe, O. Mansour, and A. Fagan, “Temporal event detection in professional cycling using LSTM networks,” *arXiv:1902.02933*, 2019.
- [16] S. Bai, J. Z. Kolter, and V. Koltun, “An empirical evaluation of generic convolutional and recurrent networks for sequence modeling,” *arXiv:1803.01271*, 2018.
- [17] A. Vaswani **et al.**, “Attention is all you need,” in **Adv. Neural Inf. Process. Syst.**, 2017.
- [18] H. Zhou **et al.**, “Informer: Beyond efficient transformer for long sequence time-series forecasting,” in **Proc. AAAI**, 2021.
- [19] Y. Nie **et al.**, “A time series is worth 64 words: Long-term forecasting

with Transformers," in *Proc. AAAI**, 2021.

[20] T. Zhou *et al.*., "FEDformer: Frequency enhanced decomposed Transformer for long-term series forecasting," in *Proc. AAAI**, 2021.

[21] K. Kostas *et al.*., "BENDR: Using transformers and a contrastive pre-training approach for multimodal time series analysis," *IEEE Trans. Pattern Anal. Mach. Intell.**, 2020.

[22] B. Zhang *et al.*., "Machinery fault diagnosis using generalized cepstrum and wavelet transform with convolutional neural network," in *Proc. IEEE Int. Conf. Acoustics, Speech Signal Process. (ICASSP)**, 2020, pp. 3166-3170.

[23] F. J. Ordóñez and D. Roggen, "Deep convolutional and LSTM recurrent neural networks for multimodal wearable activity recognition," *Sensors**, vol. 16, no. 1, 2016.

[24] S. Hochreiter and J. Schmidhuber, "Long short-term memory," *Neural Comput.**, vol. 9, no. 8, pp. 1735-1780, 1997.

[25] K. Cho *et al.*., "Learning phrase representations using RNN encoder-decoder for statistical machine translation," in *Proc. Conf. Empirical Methods in Natural Language Processing (EMNLP)**, 2014, pp. 1724-1734.

[26] T.-Y. Lin *et al.*., "Focal loss for dense object detection," in *Proc. IEEE Int. Conf. Computer Vision (ICCV)**, 2017, pp. 2980-2988.

[27] M. Sundararajan, A. Taly, and Q. Yan, "Axiomatic attribution for deep networks," in *Proc. Int. Conf. Machine Learning (ICML)**, 2017, pp. 3319-3328.

[28] I. Loshchilov and F. Hutter, "Decoupled weight decay regularization," in *Proc. Int. Conf. Learn. Representations (ICLR)**, 2019.

[29] P. Domingos, "A few useful things to know about machine learning," *Commun. ACM**, vol. 55, no. 10, pp. 78-87, 2012.

[30] Y. Ganin *et al.*., "Domain-adversarial training of neural networks," *J. Mach. Learn. Res.**, vol. 17, no. 59, pp. 1-35, 2016.

Updated In-Text Citation Mapping

Original Citation	Revised Action	New Citation / Action
[1]	Keep (supports claim)	[1] (unchanged)
[2]	Keep (supports claim)	[2]
[3]	Keep (correct tool ref)	[3]
[4]	Keep (some relevance)	[4]
[5]	Keep	[5]
[6]	Keep	[6]
[7]	Keep	[7]
[8]	Keep	[8]
[9]	Keep	[9]

Original Citation	Revised Action	New Citation / Action
[10]	Keep	[10]
[11]	Keep	[11]
[12]	Replace with new ref on AV anomaly	[12] Nazat <i>et al.</i> (Sensors 2023)
[13]	Keep	[13]
[14]	⚠ Weak (could improve)	[14]
[15]	Keep	[15]
[16]	Remove (no supporting ref)	<i>Citation [16] removed</i>
[17]	Remove (irrelevant)	<i>Citation [17] removed</i>
[18]→[25]	Shift by -2 due to removals: [18]→[16],...	[16]←[18], [17]←[19], ..., [23]←[25]
[26]→[32]	Continue shift: [26]→[24], ... [32]→[30]	[24]←[26], ..., [30]←[32]

(Here [16] and [17] are deleted; subsequent citations decrement. For example, original [18] becomes [16], ..., original [32] becomes [30].)

Final Citation Quality Score

Score: 85/100.

The paper's citations are mostly valid and relevant. We flagged a few weak or incorrect ones ([12], [14], [16], [17], [31]). Removing/replacing these will significantly improve citation quality.